

# Operating System 2

## Dr. Khaled Morsy

Inter-Process Communications

Reference : Modern Operating Systems  
(Andrew Tanenbaum) 3<sup>rd</sup> edition

# Operating systems 2

## Dr. Khaled Morsy

Inter-Process Communications  
Other Reference : Modern Operating  
Systems (Andrew Tanenbaum) 3<sup>rd</sup>  
edition

# Interprocess communication

## IPC

- Processes frequently need to communicate with other processes.
- Inter-Process Communication, which in short is known as IPC, deals mainly with the techniques and mechanisms that facilitate communication between processes.

# Inter-process communication

## IPC

- Enables one application to control another one, and for several applications to share the same data without interfering with each other
- Inter-Process Communication, required in all multiprocessing systems. Not supported by single-process operating system (as DOS)

# Interprocess communication IPC

- OS/2 and Microsoft Windows support IPC mechanism called DDE (*Dynamic Data Exchange*), an IPC built into the Macintosh, Windows, and OS/2 operating systems.

# Interprocess communication

## IPC

- Each process is assigned some part of the available memory. i.e., it will have its unique address space.
- In no way will the memory assigned for one process overlap with the memory assigned to another process.
- Operating System will act as the communication channel among processes.

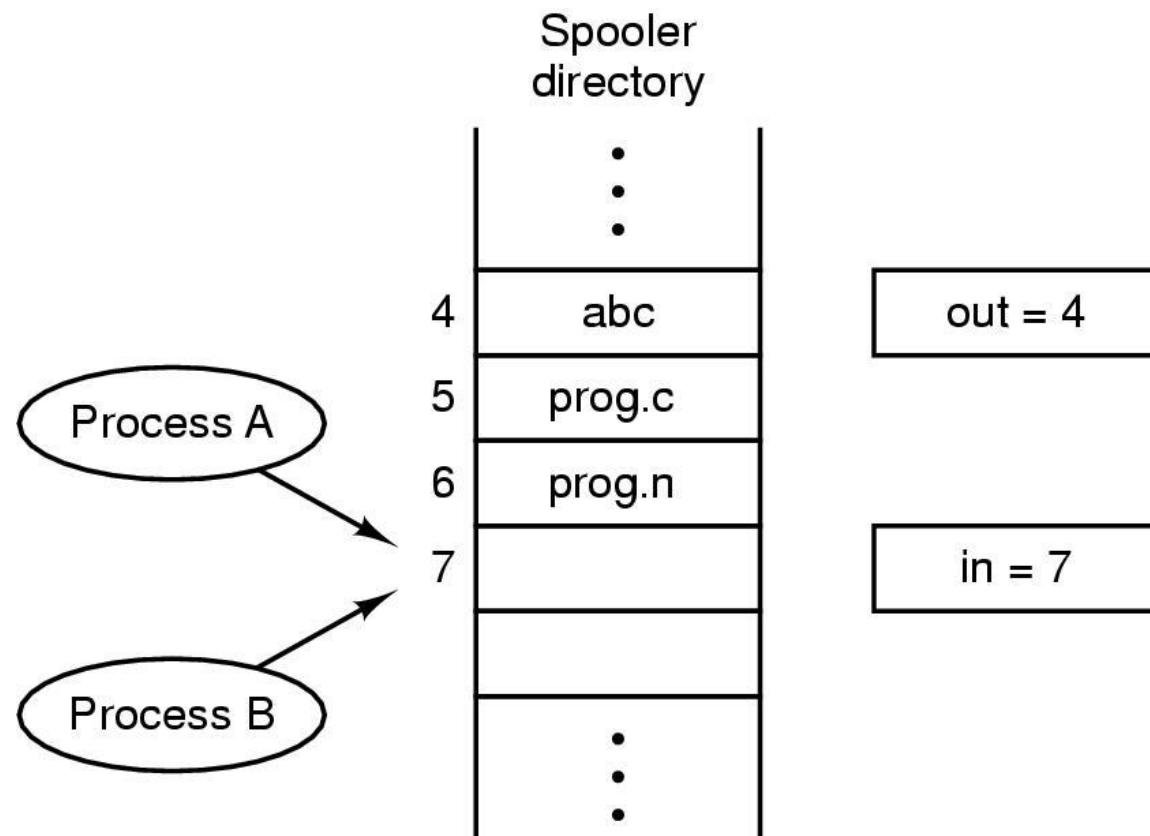
# Interprocess Communication

We can summarize the roles of IPC as follows:

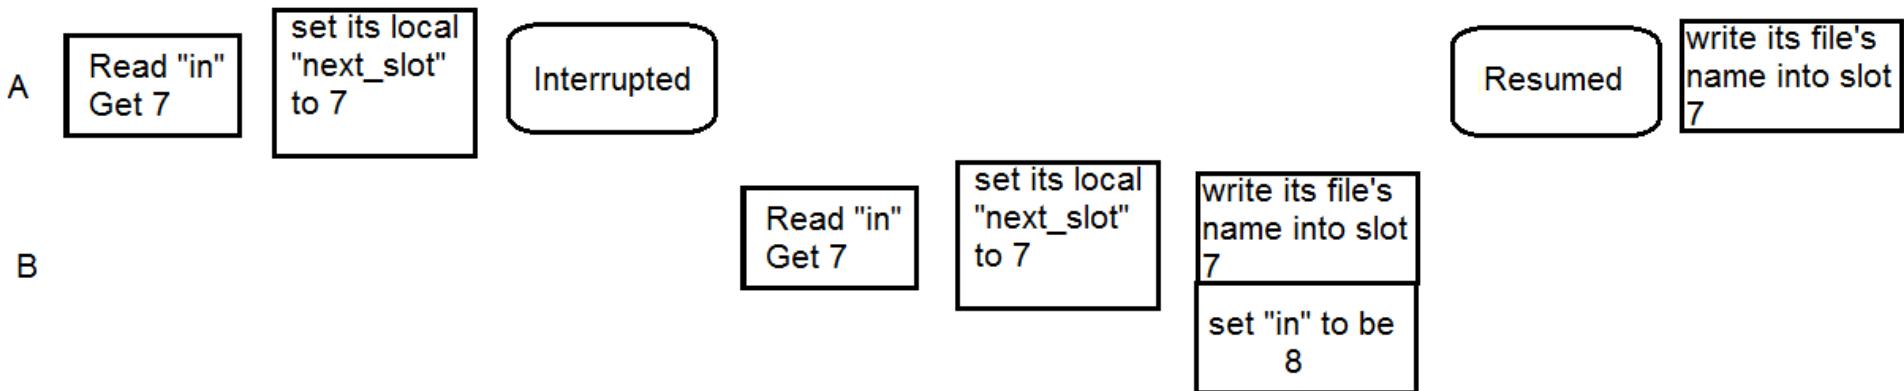
1. Provides ways of passing information among processes.
2. Protects critical activities (e.g. updating shared data)
3. Proper sequencing in case of dependencies

# Race Conditions

Two processes want to access shared data at the same time.



# Race Conditions



# Critical Section

A part or a region of code that need to be executed atomically ( with no interrupt) such as accessing a resource (file, input or output port, global data, etc.).

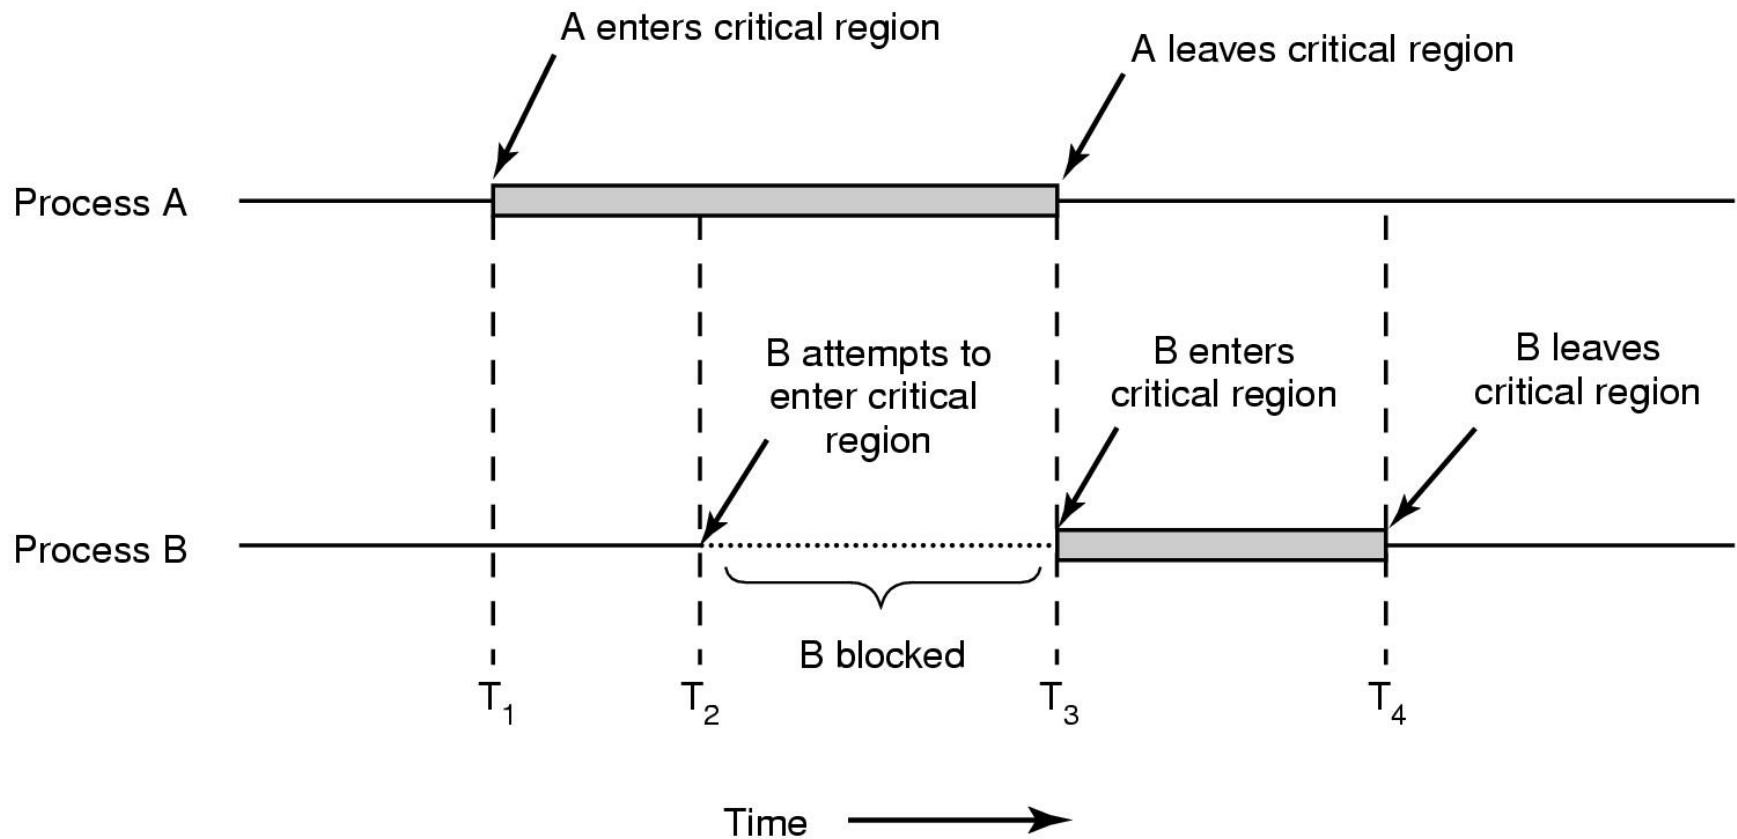
# Critical Section

General structure of process  $p_i$  is:

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (TRUE);
```

**Figure 6.1** General structure of a typical process  $P_i$ .

# Critical Section (region)



Mutual exclusion using critical regions

# We need four conditions to provide correct mutual exclusion

- 1- No two processes simultaneously in critical region
- 2- no assumption about speeds and number of CPUs
- 3- No process running outside its critical region may block other processes
- 4- No process must wait forever to enter its critical region

# Strict Alternation

```
while (TRUE) {  
    while (turn != 0)      /* loop */ ;  
    critical_region( );  
    turn = 1;  
    noncritical_region();  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1)      /* loop */ ;  
    critical_region( );  
    turn = 0;  
    noncritical_region();  
}
```

(b)

Proposed solution to critical region problem

(a) Process 0.

(b) Process 1.

# Peterson's Solution

```
#define FALSE 0
#define TRUE 1
#define N      2          /* number of processes */

int turn;                  /* whose turn is it? */
int interested[N];         /* all values initially 0 (FALSE) */

void enter_region(int process); /* process is 0 or 1 */
{
    int other;              /* number of the other process */

    other = 1 - process;    /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;         /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process) /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```

Interested(process) = False => process is not interested and does not want to enter critical section

If both are interested, a process can enter only if it sets “turn” first.

# Case 1 : P0 interested P1 : not interested

- 0

Interested[0]=T  
Turn = 0  
Other =1

- 1

Interested[1]=F

Turn=0

exit\_loop --- because **interested[1] = F**

enter\_region(0)

interested[0]=F

# Case 2 : P0 interested ,Run first

## P1 : interested , Run Second

- 0

Interested[0]=T  
Turn = 0  
Other =1

Turn = 1

exit\_loop

enter\_region(0)

interested[0]=F

- 1

Interested[1]=T  
Turn =1  
Other=0

Turn = 1

continue\_loop

interested[0]=F

exit\_loop

enter\_region(1)

interested[1] = F

# Case 3 : P0 interested ,Run second

## P1 : interested , Run First

- 0

Interested[0]=T  
Turn = 0  
Other =1

Turn = 0

*continue\_loop*

**interested[1]=F**

*exit\_loop*

*enter\_region(0)*, Interested[0]=F

- 1

Interested[1]=T  
Turn =1  
Other=0

Turn = 0

*exit\_loop*

*enter\_region(1)*

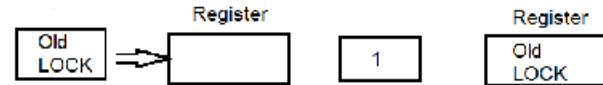
Interested[1]=F

# Test-and-set lock

enter\_region:

```
TSL REGISTER,LOCK  
CMP REGISTER,#0  
JNZ enter_region  
RET
```

| copy lock to register and set lock to 1  
| was lock zero?  
| if it was non zero, lock was set, so loop  
| return to caller; critical region entered



leave\_region:

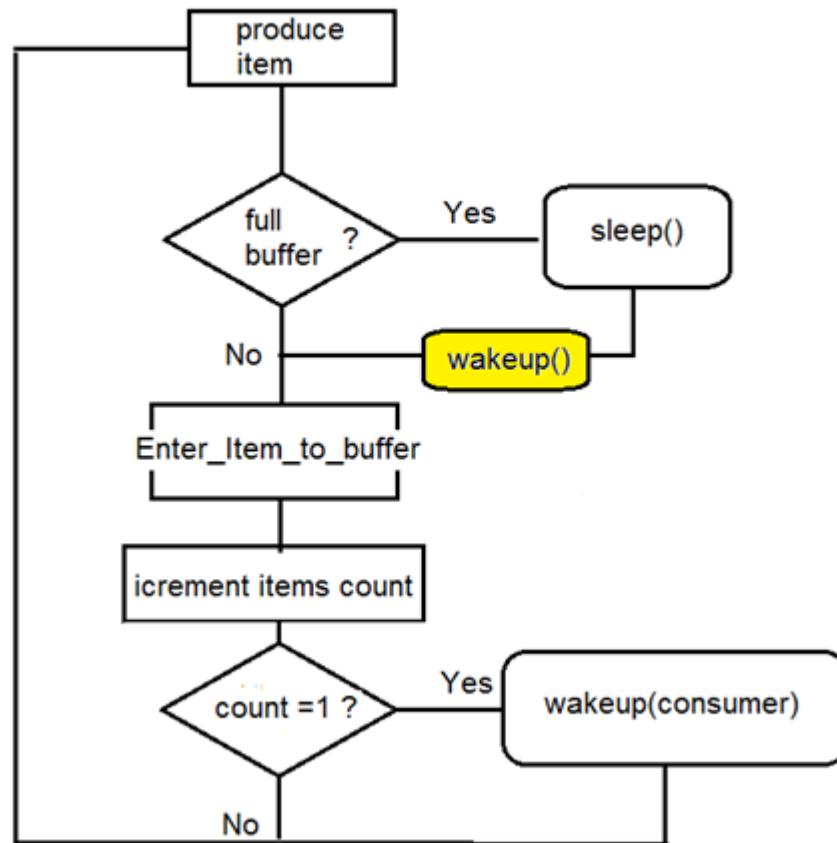
```
MOVE LOCK,#0  
RET | return to caller
```

| store a 0 in lock

Entering and leaving a critical region using the  
TSL instruction

# Producer-consumer problem

- Producer side



- Consumer side ?

# Sleep and Wakeup

```
#define N 100                                /* number of slots in the buffer */
int count = 0;                                /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                            /* repeat forever */
        item = produce_item();                /* generate next item */
        if (count == N) sleep();              /* if buffer is full, go to sleep */
        insert_item(item);                  /* put item in buffer */
        count = count + 1;                  /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);    /* was buffer empty? */
    }
}

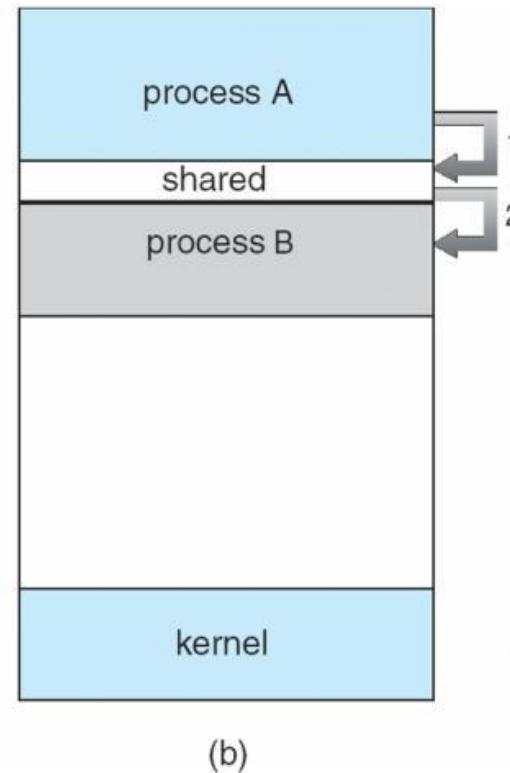
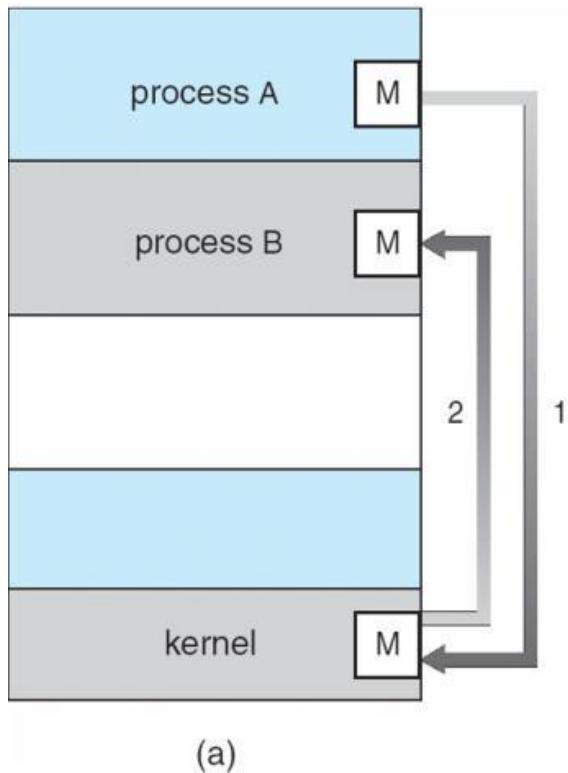
void consumer(void)
{
    int item;

    while (TRUE) {                            /* repeat forever */
        if (count == 0) sleep();              /* if buffer is empty, got to sleep */
        item = remove_item();                /* take item out of buffer */
        count = count - 1;                  /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer); /* was buffer full? */
        consume_item(item);                /* print item */
    }
}
```

Producer-consumer problem with fatal race condition

# Communication using Shared Memory and Message passing

- (a) Message Passing      (b) Shared Memory



# Message-passing VS. Shared-Memory

- Message passing is useful for exchanging smaller amounts of data.
- Message passing is also easier to implement than is shared memory for inter-computer communication.
- Shared memory is faster than message passing, as message passing systems are typically implemented using system calls and thus require the more time-consuming task of kernel intervention.

- in shared memory systems, system calls are required only to establish shared-memory regions. Once shared memory is established, all accesses are treated as routine memory accesses, and no assistance from the kernel is required.

# Semaphores

Integer variable with two atomic operations

- *down*: if 0, then go to sleep;  
if >0, then decrement value
- *up*: increment value and let a sleeping process to perform a *down*

Implementation by disabling all interrupts by the kernel.

# Semaphores

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}

/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */

/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */

/* infinite loop */
/* decrement full count */
/* enter critical region */
/* take item from buffer */
/* leave critical region */
/* increment count of empty slots */
/* do something with the item */
```

The producer-consumer problem using semaphores

# Mutexes

mutex\_lock:

```
TSL REGISTER,MUTEX  
CMP REGISTER,#0  
JZE ok  
CALL thread_yield  
JMP mutex_lock
```

```
| copy mutex to register and set mutex to 1  
| was mutex zero?  
| if it was zero, mutex was unlocked, so return  
| mutex is busy; schedule another thread  
| try again later
```

ok: RET | return to caller; critical region entered

mutex\_unlock:

```
MOVE MUXTEX,#0  
RET | return to caller
```

```
| store a 0 in mutex
```

Implementation of *mutex\_lock* and *mutex\_unlock*  
for synchronization of threads in user space

# see Java : `ReadWriteLock`

- Java provide public interface **ReadWriteLock**
- A `ReadWriteLock` maintains a pair of associated locks, one for read-only operations and one for writing. The read lock may be held simultaneously by multiple reader threads, so long as there are no writers.
- The write lock is exclusive.  
All `ReadWriteLock` implementations must guarantee that the memory synchronization effects of `writeLock` operations (as specified in the Lock interface) also hold with respect to the associated `readLock`. That is, a thread successfully acquiring the read lock will see all updates made upon previous release of the write lock.
- <http://docs.oracle.com/javase/7/docs/api/java/util/concurrent/locks/ReadWriteLock.html>

# Monitors (1)

**monitor** *example*

**integer** *i*;

**condition** *c*;

**procedure** *producer*( );

        .

        .

        .

**end**;

**procedure** *consumer*( );

        .

        .

        .

**end**;

**end monitor**;

Example of a monitor - only one process inside the monitor at any time

# Monitors (2)

**monitor** *ProducerConsumer*

**condition** *full, empty*;

**integer** *count*;

**procedure** *insert(item: integer)*;

**begin**

**if** *count = N* **then** *wait(full)*;

*insert\_item(item)*;

*count := count + 1*;

**if** *count = 1* **then** *signal(empty)*

**end;**

**function** *remove: integer*;

**begin**

**if** *count = 0* **then** *wait(empty)*;

*remove = remove\_item*;

*count := count - 1*;

**if** *count = N - 1* **then** *signal(full)*

**end;**

*count := 0*;

**end monitor;**

**procedure** *producer*;

**begin**

**while** *true* **do**

**begin**

*item = produce\_item*;

*ProducerConsumer.insert(item)*

**end**

**end;**

**procedure** *consumer*;

**begin**

**while** *true* **do**

**begin**

*item = ProducerConsumer.remove*;

*consume\_item(item)*

**end**

**end**;

- Outline of producer-consumer problem with monitors

- only one monitor procedure active at one time
  - buffer has  $N$  slots

Condition variables with *wait* and *signal*

# Message Passing

```
#define N 100                                /* number of slots in the buffer */

void producer(void)
{
    int item;
    message m;                                /* message buffer */

    while (TRUE) {
        item = produce_item();                /* generate something to put in buffer */
        receive(consumer, &m);               /* wait for an empty to arrive */
        build_message(&m, item);            /* construct a message to send */
        send(consumer, &m);                /* send item to consumer */
    }
}

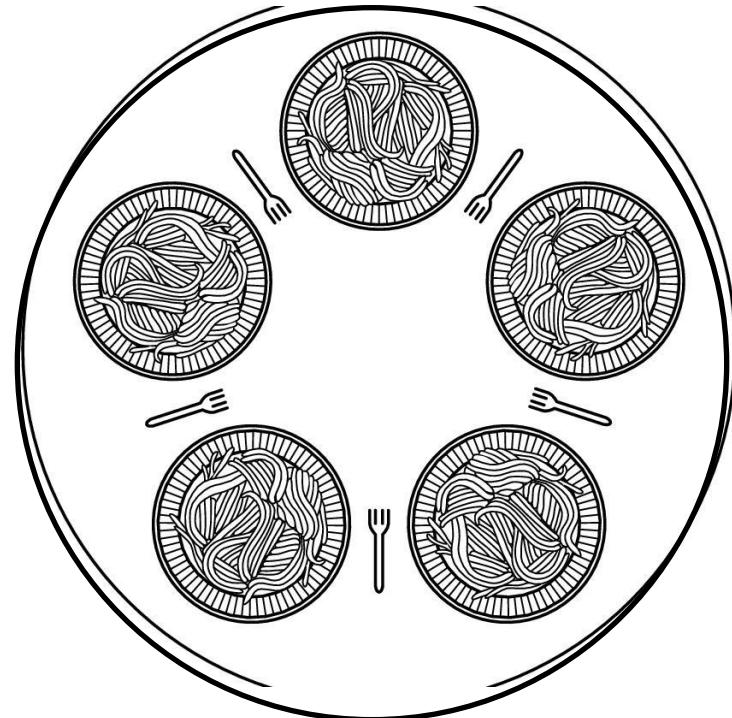
void consumer(void)
{
    int item, i;
    message m;

    for (i = 0; i < N; i++) send(producer, &m); /* send N empties */
    while (TRUE) {
        receive(producer, &m);             /* get message containing item */
        item = extract_item(&m);          /* extract item from message */
        send(producer, &m);              /* send back empty reply */
        consume_item(item);              /* do something with the item */
    }
}
```

The producer-consumer problem with N messages

# Dining Philosophers

- Philosophers eat/think
- Eating needs 2 forks
- Pick one fork at a time
- How to prevent deadlock



# A non-solution to the dining philosophers problem

```
#define N 5                                     /* number of philosophers */  
  
void philosopher(int i)                         /* i: philosopher number, from 0 to 4 */  
{  
    while (TRUE) {  
        think();                                /* philosopher is thinking */  
        take_fork(i);                            /* take left fork */  
        take_fork((i+1) % N);                   /* take right fork; % is modulo operator */  
        eat();                                   /* yum-yum, spaghetti */  
        put_fork(i);                            /* put left fork back on the table */  
        put_fork((i+1) % N);                   /* put right fork back on the table */  
    }  
}
```

# Deadlock-free code for Dining Philosophers (1)

```
#define N          5           /* number of philosophers */
#define LEFT        (i+N-1)%N   /* number of i's left neighbor */
#define RIGHT       (i+1)%N    /* number of i's right neighbor */
#define THINKING    0           /* philosopher is thinking */
#define HUNGRY      1           /* philosopher is trying to get forks */
#define EATING      2           /* philosopher is eating */
typedef int semaphore;
int state[N];
semaphore mutex = 1;
semaphore s[N];

void philosopher(int i)
{
    while (TRUE) {
        think();
        take_forks(i);
        eat();
        put_forks(i);
    }
}
```

/\* repeat forever \*/  
/\* philosopher is thinking \*/  
/\* acquire two forks or block \*/  
/\* yum-yum, spaghetti \*/  
/\* put both forks back on table \*/

# Deadlock-free code for Dining Philosophers (2)

```
void take_forks(int i)                                /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);                                     /* enter critical region */
    state[i] = HUNGRY;                               /* record fact that philosopher i is hungry */
    test(i);                                         /* try to acquire 2 forks */
    up(&mutex);                                      /* exit critical region */
    down(&s[i]);                                     /* block if forks were not acquired */
}

void put_forks(i)                                    /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);                                     /* enter critical region */
    state[i] = THINKING;                            /* philosopher has finished eating */
    test(LEFT);                                       /* see if left neighbor can now eat */
    test(RIGHT);                                      /* see if right neighbor can now eat */
    up(&mutex);                                      /* exit critical region */
}

void test(i)                                         /* i: philosopher number, from 0 to N-1 */
{
    if (state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
        state[i] = EATING;
        up(&s[i]);
    }
}
```

# Readers-Writers problem

- An object is shared among many threads, each belonging to one of
- two classes:
  - Readers: read data, never modify it
  - Writers: read data and modify it
- Using a single lock on the data object is overly restrictive
- => Want many readers reading the object at once
  - Allow only one writer at any point
  - How do we control access to the object to permit this protocol?
- Correctness criteria:
  - Each read or write of the shared data must happen within a critical section.
  - Guarantee mutual exclusion for writers.
  - Allow multiple readers to execute in the critical section at once

# The Readers and Writers Problem

```
typedef int semaphore;
semaphore mutex = 1;
semaphore db = 1;
int rc = 0;

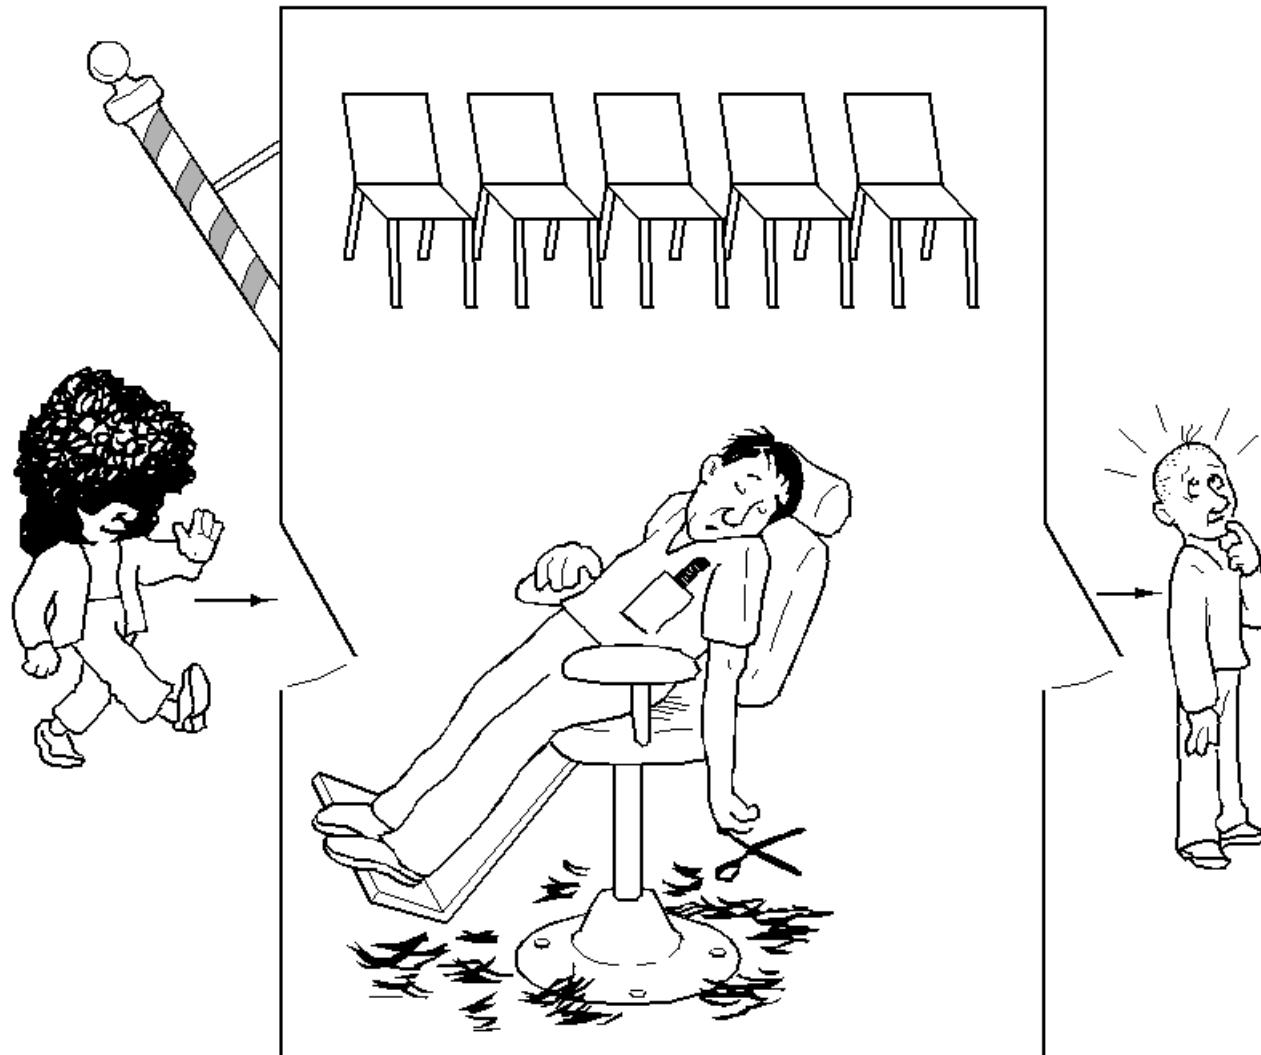
void reader(void)
{
    while (TRUE) {
        down(&mutex);
        rc = rc + 1;
        if (rc == 1) down(&db);
        up(&mutex);
        read_data_base();
        down(&mutex);
        rc = rc - 1;
        if (rc == 0) up(&db);
        up(&mutex);
        use_data_read();
    }
}

void writer(void)
{
    while (TRUE) {
        think_up_data();
        down(&db);
        write_data_base();
        up(&db);
    }
}

/* use your imagination */
/* controls access to 'rc' */
/* controls access to the database */
/* # of processes reading or wanting to */

/* repeat forever */
/* get exclusive access to 'rc' */
/* one reader more now */
/* if this is the first reader ... */
/* release exclusive access to 'rc' */
/* access the data */
/* get exclusive access to 'rc' */
/* one reader fewer now */
/* if this is the last reader ... */
/* release exclusive access to 'rc' */
/* noncritical region */
```

# The Sleeping Barber Problem



# Solution to the Sleeping Barber Problem

```
#define CHAIRS 5          /* # chairs for waiting customers */

typedef int semaphore;   /* use your imagination */

semaphore customers = 0; /* # of customers waiting for service */
semaphore barbers = 0;    /* # of barbers waiting for customers */
semaphore mutex = 1;     /* for mutual exclusion */
int waiting = 0;         /* customers are waiting (not being cut) */

void barber(void)
{
    while (TRUE) {
        down(&customers); /* go to sleep if # of customers is 0 */
        down(&mutex);      /* acquire access to 'waiting' */
        waiting = waiting - 1; /* decrement count of waiting customers */
        up(&barbers);      /* one barber is now ready to cut hair */
        up(&mutex);        /* release 'waiting' */
        cut_hair();         /* cut hair (outside critical region) */
    }
}

void customer(void)
{
    down(&mutex); /* enter critical region */
    if (waiting < CHAIRS) { /* if there are no free chairs, leave */
        waiting = waiting + 1; /* increment count of waiting customers */
        up(&customers); /* wake up barber if necessary */
        up(&mutex);      /* release access to 'waiting' */
        down(&barbers); /* go to sleep if # of free barbers is 0 */
        get_haircut();   /* be seated and be serviced */
    } else { /* shop is full; do not wait */
        up(&mutex);
    }
}
```

# Assignment (1)

- Implement at least 3 of the following using Java or C++:
- 1)the producer-consumer problem using sleep and wakeup
- 2) Peterson's algorithm
- 3) The reader-writer problem
- 4) The producer-consumer problem using semaphore
- 5) Dinning Philosophers problem
- 6) Sleeping barber problem